



Data Structures: Lecture 13

Herbert H. Chang, Ph.D.

張恆華

Eng Sci Ocean Eng

National Taiwan University

Quick Sort, Disjoint Sets, and Selection

H. Chang © NTU

Quicksort

- Quicksort is a recursive divide-and-conquer algorithm, like mergesort.
- Quicksort is the fastest known comparison-based sort for arrays, even though it has a $\Theta(n^2)$ worst-case time.
- If properly designed, it virtually always runs in $O(n \log n)$ in practice.
- Given an unsorted list I of items, quicksort chooses a "pivot" item v from I , then puts each item of I into one of two unsorted lists, depending on whether its key is less or greater than v 's key.

2

H. Chang © NTU

Algorithm

Start with unsorted list I of n items

Choose pivot item v from I

Partition I into 2 unsorted lists I_1 & I_2

I_1 : All keys smaller than v 's key.

I_2 : All keys larger than v 's key.

Items with same key as v can go into either list.

The pivot does not go into either list.

Sort I_1 recursively, yielding sorted list S_1

Sort I_2 recursively, yielding sorted list S_2


Concatenate S_1 , v , & S_2 , yielding sorted list S

3

H. Chang © NTU

Example

- The pivot v is always chosen to be the first item.




- We get lucky: pivot always be the item having the median key.
- Running time: $O(n \log n)$.

4

H. Chang © NTU

Example: already sorted list



- When the input list I happens to be already sorted, choosing the pivot to be the first item of the list is a disastrous policy.

5

H. Chang © NTU

Choosing a pivot


- We need a better way to choose a pivot.
- Randomly select an item from I to serve as pivot.
- With a random pivot, we can expect “on average” to obtain a $\frac{1}{4} \sim \frac{3}{4}$ split.
- The average running time with random pivots is $O(n \log n)$.
- An even better way to choose a pivot (when n is larger than 50 or so) is called the “median-of-three” strategy.
- “Median-of-three”: select 3 random items from I , then choose the middle key among them.

6

H. Chang © NTU

Quicksort on linked lists

- Suppose we put all items with same key as v into I_1 .
- If we try to sort a list in which every single item has the same key, then "every" item will go into list I_1 , and quicksort will have quadratic running time!.



7

H. Chang © NTU

Far better solution

- When sorting a linked list, a far better solution is to partition I into 3 unsorted lists: I_1 , I_2 , & I_v . I_v contains pivot & all items with same key as v .
- Sort I_1 & I_2 recursively, yielding S_1 and S_2 ; not I_v .
- Finally, concatenate S_1 , I_v , S_2 to yield S .



- Unfortunately, with linked lists, selecting a pivot is annoying.

8

H. Chang © NTU

Quicksort on arrays

- In-place quicksort is very fast. But a fast in-place quicksort is tricky to code.
- Suppose we have an array **a**.
- We want to sort the items starting at **a[low]** and ending at **a[high]**.
- We choose a pivot **v**; swap it with the last item, **a[high]**.




9

H. Chang © NTU

Quicksort on arrays

- We employ two array indices, **i** and **j**. Index **i** is initially "low - 1", and index **j** is initially "high".
- Indices **i** & **j** sandwich items to sort (not including pivot).
- Invariants:
 - All items at or left of index **i** have a key \leq pivot's key.
 - All items at or right of index **j** have a key \geq pivot's key.
- Repeatedly advance **i** to key \geq pivot.
- Repeatedly decrement **j** to key \leq pivot.



10

H. Chang © NTU

Quicksort on arrays

- Then, swap items at i & j .

3	4	0	9	8	7	5
		i		j		

- We repeat this sequence until the indices i and j meet in the middle, $i \geq j$.

3	4	0	9	8	7	5
		j	i			

- Then, we move the pivot back into the middle by swapping the pivot with the item at index i .

3	4	0	5	8	7	9
		$I1$		$I2$		

11

H. Chang © NTU

Quicksort on arrays

- How about items with same key as pivot?
→ Handling these is particularly tricky.
- We'd like to put them on a separate list (as we did for linked lists) → in-place is too complicated.
⇒ Make sure each index i & j , stops whenever it reaches a key equal to the pivot.
- Every key equal to the pivot takes part in one swap.
- Swapping an item equal to the pivot may seem unnecessary, but it has an excellent side effect: if all the items in the array have the same key, half these items will go into $I1$, and half into $I2$, giving us a well-balanced recursion tree.

12

H. Chang © NTU

Pseudocode

```
void quicksort(Comparable[] a, int low, int high) {  
    // Comparable is some variable type supporting  
    // comparable function  
    if(low < high) {  
        int pivotIndex = random number from low to high;  
        Comparable pivot = a[pivotIndex];  
        a[pivotIndex] = a[high];  
        // swap pivot with last item  
        a[high] = pivot;  
        int i = low - 1;  
        int j = high;
```

13

H. Chang © NTU

Pseudocode

```
do {  
    do{i++;} while (a[i].compareTo(pivot)<0);  
    do{j--;} while ((a[j].compareTo(pivot)>0)  
        && (j>low));  
    if(i < j) {  
        swap a[i] and a[j];  
    }  
} while (i < j);
```

14

H. Chang © NTU

Pseudocode

```
a[high] = a[i];  
a[i] = pivot;  
// put pivot in middle  
quicksort(a, low, i-1);  
// recursively sort left list  
quicksort(a, i+1, high);  
// recursively sort right list  
}  
}
```

15

H. Chang © NTU

Disjoint Sets


- A disjoint sets data structure represents a collection of sets that are disjoint.
- No item is in more than one set.
- A collection of disjoint sets is called a partition.
- Universe of items: all of the items that can be a member of a set. Each item is a member of exactly one set.

16

H. Chang © NTU

Two operations

- Union: merges two sets into one.
- Find: takes an item and tells us what set it is in.
- Data structure designed to support these operations are called partition or union/find data structures.
- Applications: Kruskal's algorithm for computing the minimum spanning tree of a graph.




17

H. Chang © NTU

Example: corporations

- Union/find data structures begin with every item in a separate set.
- The query "find(Empire Air)" returns "Empire Air".



- After 3 unions: "find(Empire Air)" returns "Piedmont Air".
- After the final union: "find(Empire Air)" returns "US Air".

18

H. Chang © NTU

List-Based Disjoint Sets

- Each set references a lists of the items in that set.
- Each item references the set that contains it.
- With this data structure, Find operations take $O(1)$ time.
- Union operations are slow.
- List-based disjoint sets use the quick-find algorithm.

19

H. Chang © NTU

Tree-Based Disjoint Sets

- Union operations take $O(1)$ time.
- Find operations are slower.
- Quick-union is faster overall than quick-find.
- To support fast unions, each set is maintained as a tree.
- The quick-union data structure comprises a forest (a collection of trees).

20

H. Chang © NTU

The Quick-Union Algorithm

- Each item is initially the root of its own tree.
- No child or sibling references; only parent.
- The true identity of each set is recorded at its root.
- Union: make the root of one set be a child of the root of the other set.




- Find: follow the chain of parent references from an item to the root of its tree.

21

H. Chang © NTU

Union-by-size

- The cost of Find operation is proportional to the item's depth in the tree.
- Union-by-size: keep items from getting too deep by uniting sets intelligently.
- At each root, we record the size of its tree.
- When we unite two trees, we make the smaller one a subtree of the larger one, breaking ties arbitrarily.
- e.g.



22

H. Chang © NTU

Implementing Quick-Union with an Array

- Suppose the items are non-negative integers, numbered from zero.
- We'll use an array to record the parent of each item.
- If an item has no parent, we'll record the size of its tree, record the size s as the negative number $-s$.
- Initially, every item is the root of its own tree, so we set every array element to -1 .

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
0	1	2	3	4	5	6	7	8	9

23

H. Chang © NTU

Example

- The forest illustrated below is represented by the array beneath it.



1	-4	-1	8	5	8	1	3	-5	1
0	1	2	3	4	5	6	7	8	9

24

H. Chang © NTU

Quick-Union-by-Size

- Let root1 and root2 be two items that are roots of their respective trees.


```
void union (int root1, int root2) {  
    if (array[root2] < array[root1]) {  
        // root2 has larger tree  
        array[root2] += array[root1];  
        // update # of items in root2's tree  
        array[root1] = root2;  
        // make root2 parent of root1  
    } else { // root1 has equal or larger tree  
        array[root1] += array[root2];  
        array[root2] = root1;  
    }  
}
```

25

H. Chang © NTU

Path compression

- Suppose a tall tree, and we perform find() repeatedly on its deepest leaf.
- find(): we walk up the tree from leaf to root, move every node we encounter up the tree so that it becomes a child of the root.



26

- This technique is called “path compression”

H. Chang © NTU

Find with path compression

- Let x be an item whose set we wish to identify:
- Here is code for find, which returns the identity of the item at the root of the tree.

```
int find(int x) {  
    if (array[x] < 0) { // x is root of tree; return it  
        return x;  
    } else {  
        // find out root; compress path by making root x's parent  
        array[x] = find(array[x]);  
        return array[x];    // return the root  
    }  
}
```

27

H. Chang © NTU

Running Time

- Union: $\Theta(1)$ time
- Find: $\Theta(\log u)$ worst-case time, where u is number of unions prior to the find.
- Average running time close to constant thanks to path compression.
- Bottom line: a sequence of f find and u union operations takes $\Theta(u + f\alpha(f + u, u))$ time worst case, where α is an extremely slow-growing “inverse Ackermann function”, never > 4 for any values of f & u you’ll ever use.

28

H. Chang © NTU

Selection

- Suppose that we want to find the k th smallest key in a list.
- In other words, we want to know which item has index j , where $j = k - 1$, if the list is sorted.
- But if we don't actually need to sort the list, is there a faster way?
- This problem is called "selection".
- Example: Median of a set of n keys \rightarrow the item whose index is $j = (n-1)/2$ in the sorted list, where n is odd.

29

H. Chang © NTU

Quickselect

- We can modify quicksort to perform selection.
- Observe that when we choose a pivot v and use it to partition the list into three lists I_1 , I_v , and I_2 , we know which of the three lists contains index j , because we know the lengths of I_1 and I_2 .
- Therefore, we only need to search one of the three lists.

30

H. Chang © NTU

Quickselect algorithm

- Here's the quickselect algorithm for finding the item at index j – i.e., having the $(j+1)$ th smallest key.

Start with an unsorted list I of n items.

Choose a pivot v from I .

Partition I into three lists I_1 , I_v , and I_2 .

- I_1 contains all items whose keys are smaller than v 's key.
- I_2 contains all items whose keys are larger than v 's.
- I_v contains the pivot v .
- Items with the same key as V go into any list. (List-based: I_v ; array-based: $I_1 \& I_2$).

31

H. Chang © NTU

Quickselect algorithm

```
if (j < |I1|) {  
    Recursively find the item with index j in  
    I1; return it.  
} else if (j < |I1|+|Iv|) {  
    return the pivot v.  
} else { // j >= |I1| + |Iv|  
    Recursively find the item with index j-  
    |I1|-|Iv| in I2; return it.  
}
```

32

H. Chang © NTU

Quickselect

- The advantage of quickselect over quicksort is that we only have to make one recursive call, instead of two.
- Quickselect is much faster than quicksort, it runs in $\Theta(n)$ average time if we select pivots randomly.
- We can easily modify the code for quicksort on arrays to do selection.
- Recall that when the partition stage finishes, the pivot is stored at index "i". In the quickselect pseudocode, just replace $|I_l|$ with i and $|I_v|$ with 1.


33

H. Chang © NTU

Lower Bound on Comparison-Based Sorting

- Suppose we have a scrambled array of n numbers with each from $1 \dots n$ occurring once.
- How many possible orders can they be in?
- The answer is $n!$

$$n! = 1 \times 2 \times 3 \times \dots \times (n-2) \times (n-1) \times n$$



- Each order is a permutation of the numbers, and there are $n!$ possible permutations.

34

H. Chang © NTU

Lower Bound Analysis

- Observe that if $n > 0$:

$$n! = 1 \times 2 \times 3 \times \cdots \times (n-1) \times n \leq \overbrace{n \times n \times \cdots \times n}^{n \text{ times}} = n^n$$

and (supposing n is even)

$$\begin{aligned} n! &\geq \overbrace{\frac{n}{2} \times (\frac{n}{2} + 1) \times \cdots \times (n-1) \times n}^{n/2 \text{ times}} \\ &\geq \underbrace{\frac{n}{2} \times \frac{n}{2} \times \cdots \times \frac{n}{2} \times \frac{n}{2}}_{n/2 \text{ times}} = \left(\frac{n}{2}\right)^{n/2} \end{aligned}$$

So $n!$ is between $\left(\frac{n}{2}\right)^{n/2}$ and n^n .

$$\log\left(\frac{n}{2}\right)^{n/2} = \frac{n}{2} \log \frac{n}{2} \in \Theta(n \log n) \text{ and } \log(n)^n = n \log n$$

Hence, $\log(n!) \in \Theta(n \log n)$.

35

H. Chang © NTU

Comparison-Based Sort

- A comparison-based sort is one in which all decisions are based on comparing keys (using “if” statements).
- A correct sorting algorithm must generate a different sequence of true/false answers for each different permutation of $1 \dots n$. (Because it takes a different sequence of data moments to sort each permutation.)
- There are $n!$ different permutations $\rightarrow n!$ different sequences of true/false answers.

36

H. Chang © NTU

Comparison-Based Sort

- If a sorting algorithm asks $\leq d$ true/false questions, it generates $\leq 2^d$ different sequences of true/false answers.
 - If it correctly sorts every permutation, then $n! \leq 2^d$, so $\log(n!) \leq d$ and $d \in \Omega(n \log n)$.
 - Algorithm spends $\Theta(d)$ time asking d sequences.
- EVERY comparison-based sorting algorithm takes $\Omega(n \log n)$ worst-case time.**
- Fast algorithms make q -way decisions for large q , instead of true/false (2-way) decisions.
 - Some of these algorithms run in linear time.