

Parallel Programming in C with MPI and OpenMP

Michael J. Quinn



Chapter 1

Motivation and History



Outline

- Motivation
- Modern scientific method
- Evolution of supercomputing
- Modern parallel computers
- Seeking concurrency
- Data clustering case study
- Programming parallel computers

Why Faster Computers?

- Solve compute-intensive problems faster
 - Make infeasible problems feasible
 - Reduce design time
- Solve larger problems in same amount of time
 - Improve answer's precision
 - Reduce design time
- Gain competitive advantage

Definitions

- Parallel computing
 - Using parallel computer to solve single problems faster
- Parallel computer
 - Multiple-processor system supporting parallel programming
- Parallel programming
 - Programming in a language that supports concurrency explicitly


Why MPI?

- MPI = “Message Passing Interface”
- Standard specification for message-passing libraries
- Libraries available on virtually all parallel computers
- Free libraries also available for networks of workstations or commodity clusters

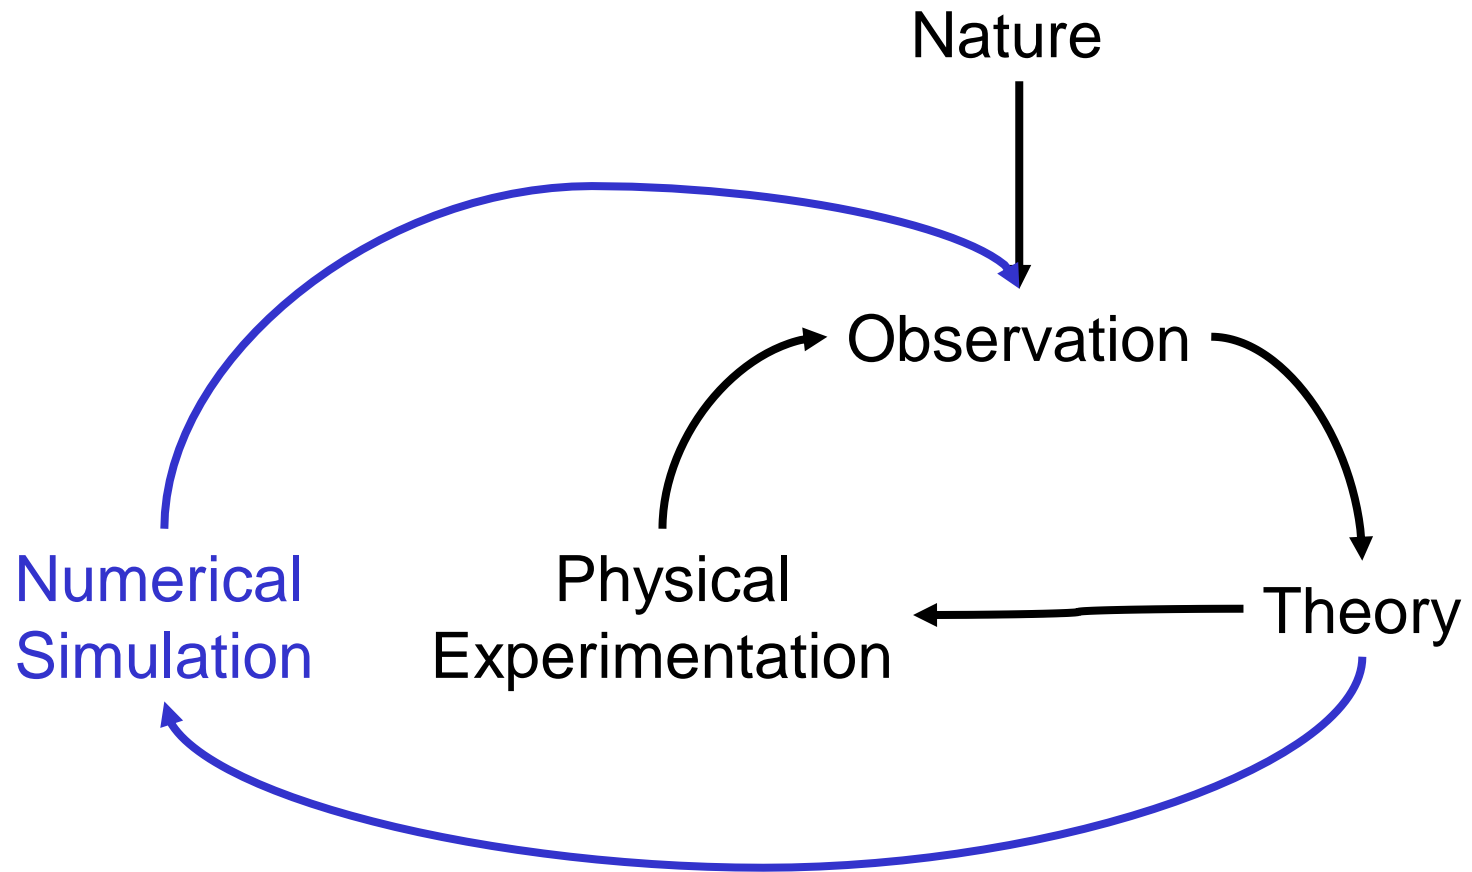
Why OpenMP?

- OpenMP an application programming interface (API) for shared-memory systems
- Supports higher performance parallel programming of symmetrical multiprocessors

Classical Science



Modern Scientific Method



Evolution of Supercomputing

- World War II
 - Hand-computed artillery tables
 - Need to speed computations
 - ENIAC
- Cold War
 - Nuclear weapon design
 - Intelligence gathering
 - Code-breaking


Supercomputer

- General-purpose computer
- Solves individual problems at high speeds, compared with contemporary systems
- Typically costs \$10 million or more
- Traditionally found in government labs

Commercial Supercomputing

- Started in capital-intensive industries
 - Petroleum exploration
 - Automobile manufacturing
- Other companies followed suit
 - Pharmaceutical design
 - Consumer products

50 Years of Speed Increases




CPU's 1 Million Times Faster


- Faster clock speeds
- Greater system concurrency
 - Multiple functional units
 - Concurrent instruction execution
 - Speculative instruction execution

Systems 1 Billion Times Faster


- Processors are 1 million times faster
- Combine thousands of processors
- Parallel computer
 - Multiple processors
 - Supports parallel programming
- Parallel computing = Using a parallel computer to execute a program faster

Microprocessor Revolution





<http://www.acmqueue.org/figures/issue007/borkarfig1.gif>



<http://www.cochrane.org.uk/opinion/papers/images/exponentialeducation4.gif>

Evolution of Computer Power/Cost

MIPS per \$1000 (1997 Dollars)






Figure E1A


Moore's Law: The Growth of Computing, 1900-1998



Source: Bill Gates *The Road Ahead* (New York: Viking, 1995) p. 32; Ray Kurzweil *The Age of Spiritual Machines* (New York: Penguin, 1999) pp. 22-4.



<http://www.intel.com/research/silicon/images/MooresLawgraph3.gif>



<http://www.nature.com/nbt/journal/v21/n10/images/nbt871-F4.gif>

Modern Parallel Computers

- Caltech's Cosmic Cube (Seitz and Fox)
- Commercial copy-cats
 - nCUBE Corporation
 - Intel's Supercomputer Systems Division
 - Lots more
- Thinking Machines Corporation

Copy-cat Strategy

- Microprocessor
 - 1% speed of supercomputer
 - 0.1% cost of supercomputer
- Parallel computer = 1000 microprocessors
 - 10 x speed of supercomputer
 - Same cost as supercomputer

Why Didn't Everybody Buy One?

- Supercomputer $\neq \Sigma$ CPUs
 - Computation rate \neq throughput
 - Inadequate I/O
- Software
 - Inadequate operating systems
 - Inadequate programming environments

After the “Shake Out”

- IBM
- Hewlett-Packard
- Silicon Graphics
- Sun Microsystems

Commercial Parallel Systems

- Relatively costly per processor
- Primitive programming environments
- Focus on commercial sales
- Scientists looked for alternative

Beowulf Concept

- NASA (Sterling and Becker)
- Commodity processors
- Commodity interconnect
- Linux operating system
- Message Passing Interface (MPI) library
- High performance/\$ for certain applications

Advanced Strategic Computing Initiative

- U.S. nuclear policy changes
 - Moratorium on testing
 - Production of new weapons halted
- Numerical simulations needed to maintain existing stockpile
- Five supercomputers costing up to \$100 million each

ASCI White (10 teraops/sec)

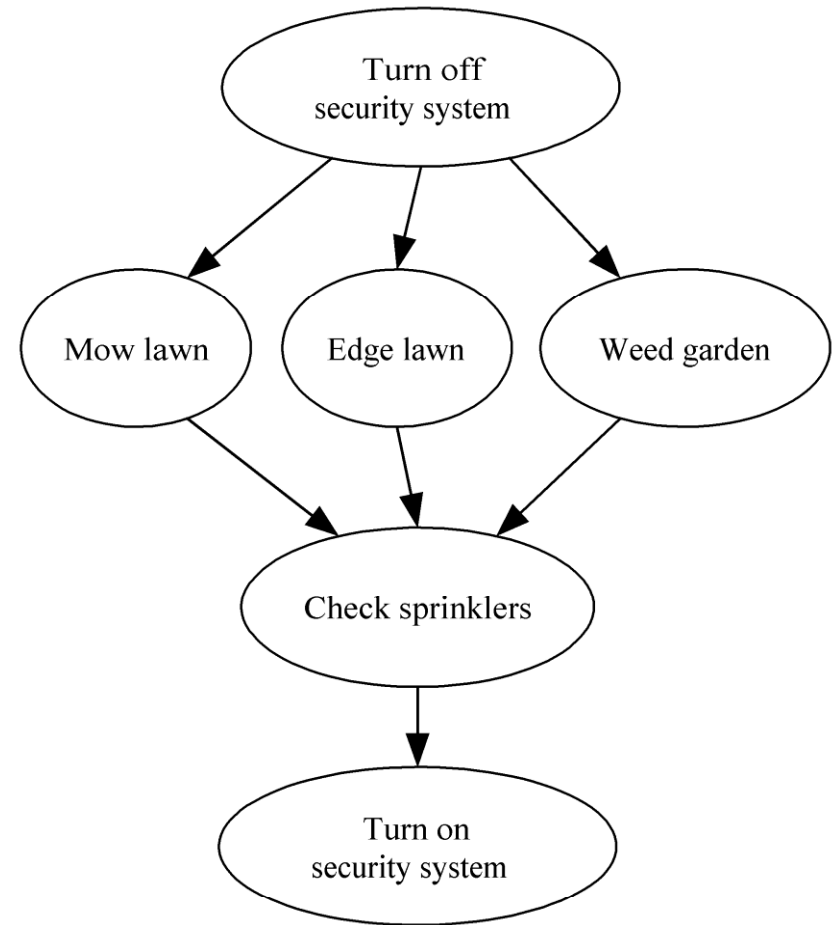


Seeking Concurrency

- Data dependence graphs
- Data parallelism
- Functional parallelism
- Pipelining

Data Dependence Graph

- Directed graph
- Vertices = tasks
- Edges = dependences



Data Parallelism

- Independent tasks apply same operation to different elements of a data set

```
for i ← 0 to 99 do  
    a[i] ← b[i] + c[i]  
endfor
```

- Okay to perform operations concurrently

Functional Parallelism

- Independent tasks apply different operations to different data elements

$a \leftarrow 2$

$b \leftarrow 3$

$m \leftarrow (a + b) / 2$


$s \leftarrow (a^2 + b^2) / 2$

$v \leftarrow s - m^2$

- First and second statements
- Third and fourth statements


Pipelining

- Divide a process into stages
- Produce several items simultaneously



Partial Sums Pipeline

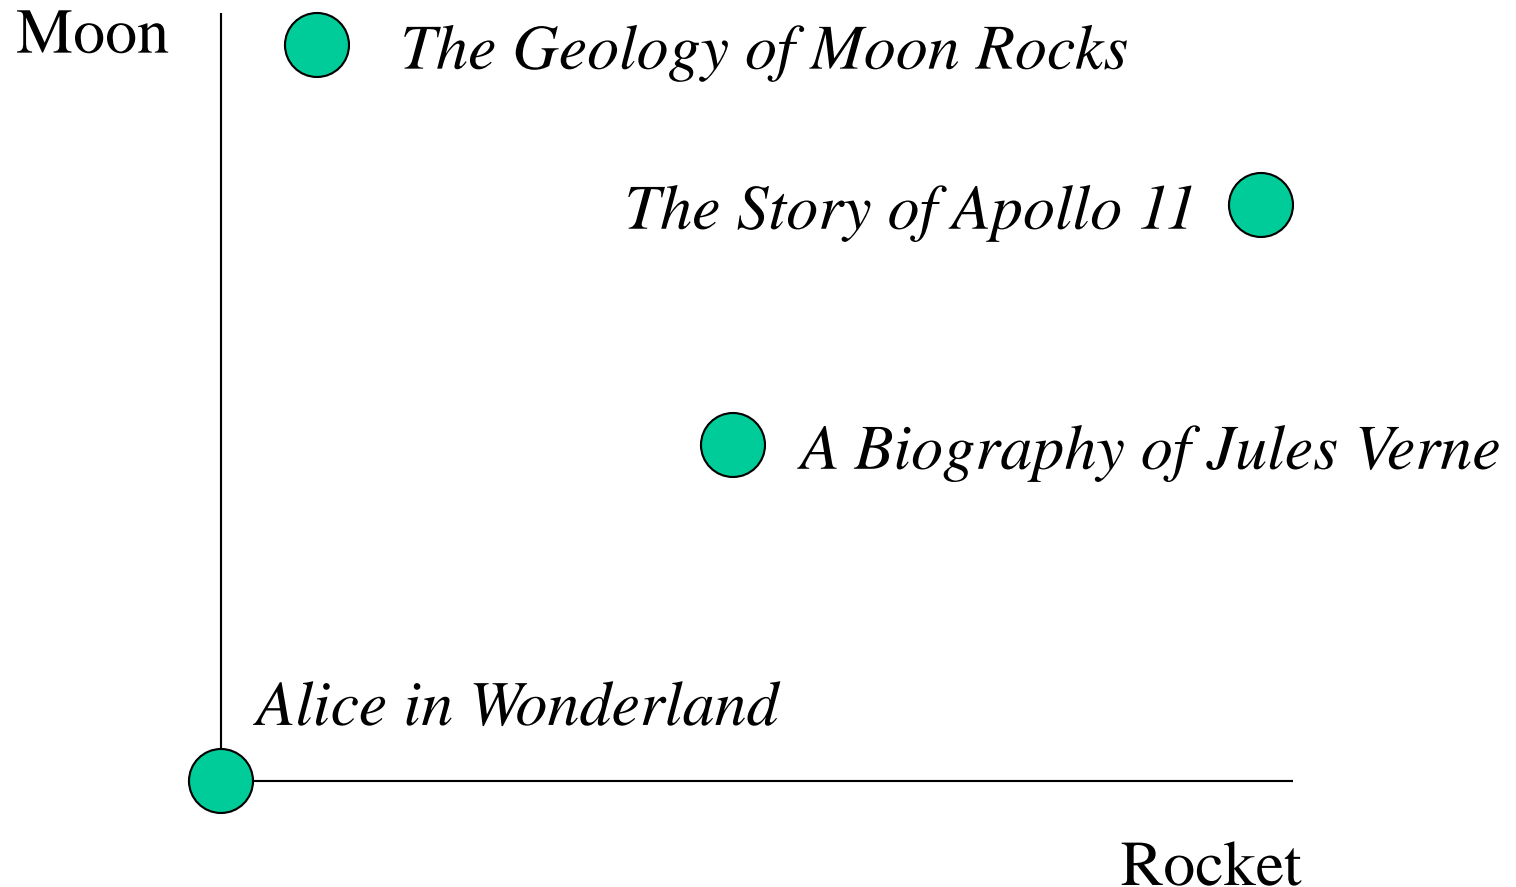
```
p[0] ← a[0]
for i ← 1 to 3 do
    p[i] ← p[i-1] + a[i]
endfor
```




Data Clustering

- Data mining
 - looking for meaningful patterns in large data sets
- Data clustering
 - organizing a data set into clusters of “similar” items
- Data clustering can speed retrieval of related items

Document Vectors



Document Clustering



Clustering Algorithm

- Compute document vectors
- Choose initial cluster centers
- Repeat
 - Compute performance function
 - Adjust centers
- Until function value converges or max iterations have elapsed
- Output cluster centers

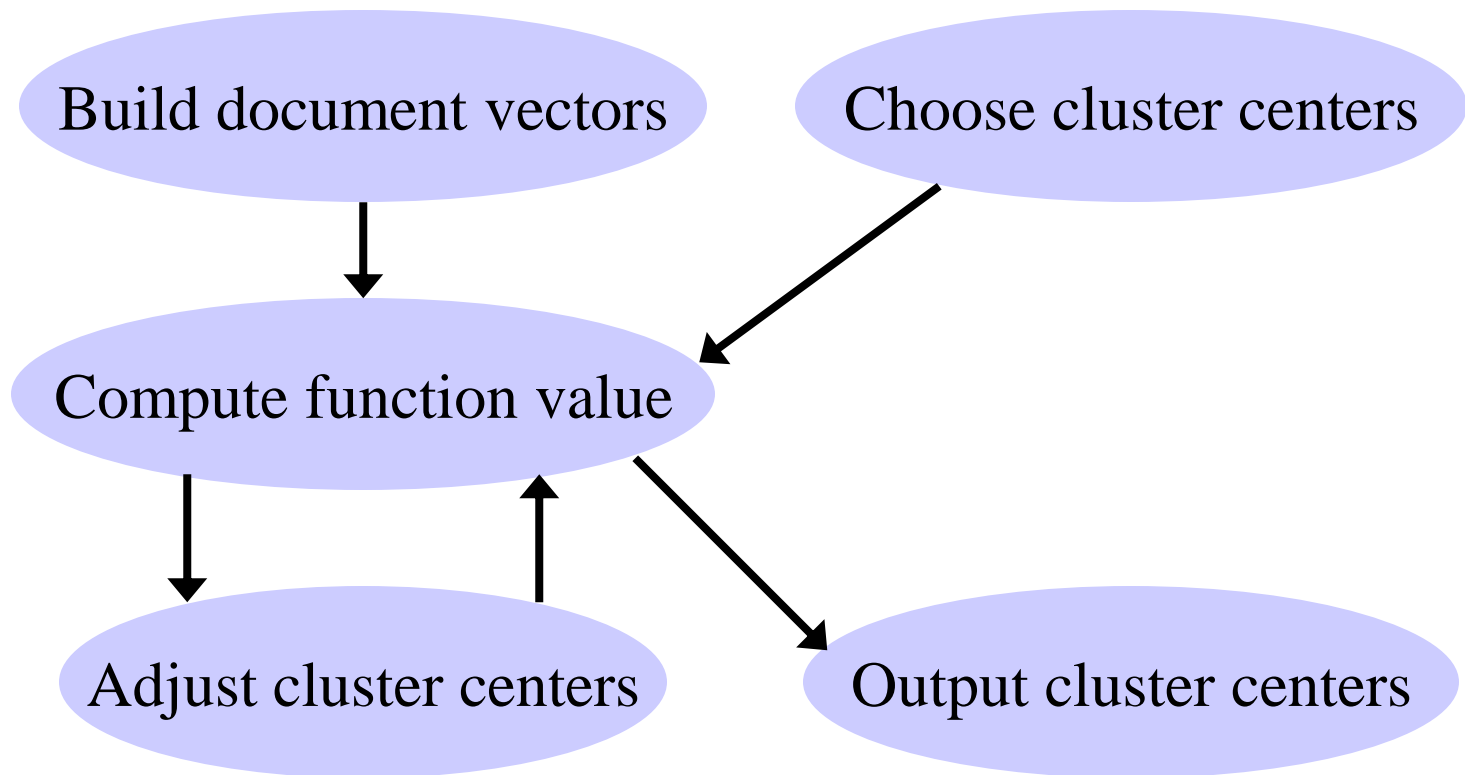
Data Parallelism Opportunities

- Operation being applied to a data set
- Examples
 - Generating document vectors
 - Finding closest center to each vector
 - Picking initial values of cluster centers

Functional Parallelism Opportunities

- Draw data dependence diagram
- Look for sets of nodes such that there are no paths from one node to another

Data Dependence Diagram



Programming Parallel Computers

- **Extend compilers**: translate sequential programs into parallel programs
- **Extend languages**: add parallel operations
- **Add parallel language layer** on top of sequential language
- Define totally **new parallel language** and compiler system

1. Extend Compilers

- Parallelizing compiler
 - Detect parallelism in sequential program
 - Produce parallel executable program
- Focus on making Fortran programs parallel

1. Extend Compilers (cont.)

- Advantages
 - Can leverage millions of lines of existing serial programs
 - Saves time and labor
 - Requires no retraining of programmers
 - Sequential programming easier than parallel programming
- Disadvantages
 - Parallelism may be irretrievably lost when programs written in sequential languages
 - Performance of parallelizing compilers on broad range of applications still up in air

2. Extend Language (Using API)

- Add functions to a sequential language
 - Create and terminate processes
 - Synchronize processes
 - Allow processes to communicate
- Example
 - PVM, MPI standard, OpenMP

2. Extend Language (cont.)

- Advantages
 - Easiest, quickest, and least expensive
 - Allows existing compiler technology to be leveraged
 - New libraries can be ready soon after new parallel computers are available
- Disadvantages
 - Lack of compiler support to catch errors
 - Easy to write programs that are difficult to debug

3. Add a Parallel Programming Layer

- Lower layer
 - Core of computation
 - Process manipulates its portion of data to produce its portion of result
- Upper layer
 - Creation and synchronization of processes
 - Partitioning of data among processes
- A few research prototypes have been built based on these principles – CODE, Hence

4. Create a Parallel Language

- Develop a parallel language “from scratch”
 - Erlang, occam is an example
- Add parallel constructs to an existing language
 - Fortran 90
 - High Performance Fortran
 - C*

4. New Parallel Languages (cont.)

- Advantages
 - Allows programmer to communicate parallelism to compiler
 - Improves probability that executable will achieve high performance
- Disadvantages
 - Requires development of new compilers
 - New languages may not become standards
 - Programmer resistance

Current Status

- Low-level approach is most popular
 - Augment existing language with low-level parallel constructs
 - MPI and OpenMP are examples
- Advantages of low-level approach
 - Efficiency
 - Portability
- Disadvantage: More difficult to program and debug

Summary (1/2)

- High performance computing
 - U.S. government
 - Capital-intensive industries
 - Many companies and research labs
- Parallel computers
 - Commercial systems
 - Commodity-based systems

Summary (2/2)

- Power of CPUs keeps growing exponentially
- Parallel programming environments changing very slowly
- Two standards have emerged
 - MPI library, for processes that do not share memory
 - OpenMP directives, for processes that do share memory

Exercise

- 1.10
- 1.13