

Parallel Programming in C with MPI and OpenMP

Michael J. Quinn



Chapter 3

Parallel Algorithm Design




Outline

- Task/Channel Model
- Algorithm Design Methodology
- Case Studies
 - Boundary value problem
 - Finding the maximum
 - The n-body problem
 - Adding data input

Task/Channel Model

- Parallel computation
 - a set of **tasks** interact by sending messages through **channels**
- Task
 - Program + Local memory + Collection of I/O ports
- Channel
 - a message queue (output port → input port)
 - Receiving is a **synchronous** operation (wait → block)
Sending is an **asynchronous** operation

Task/Channel Model



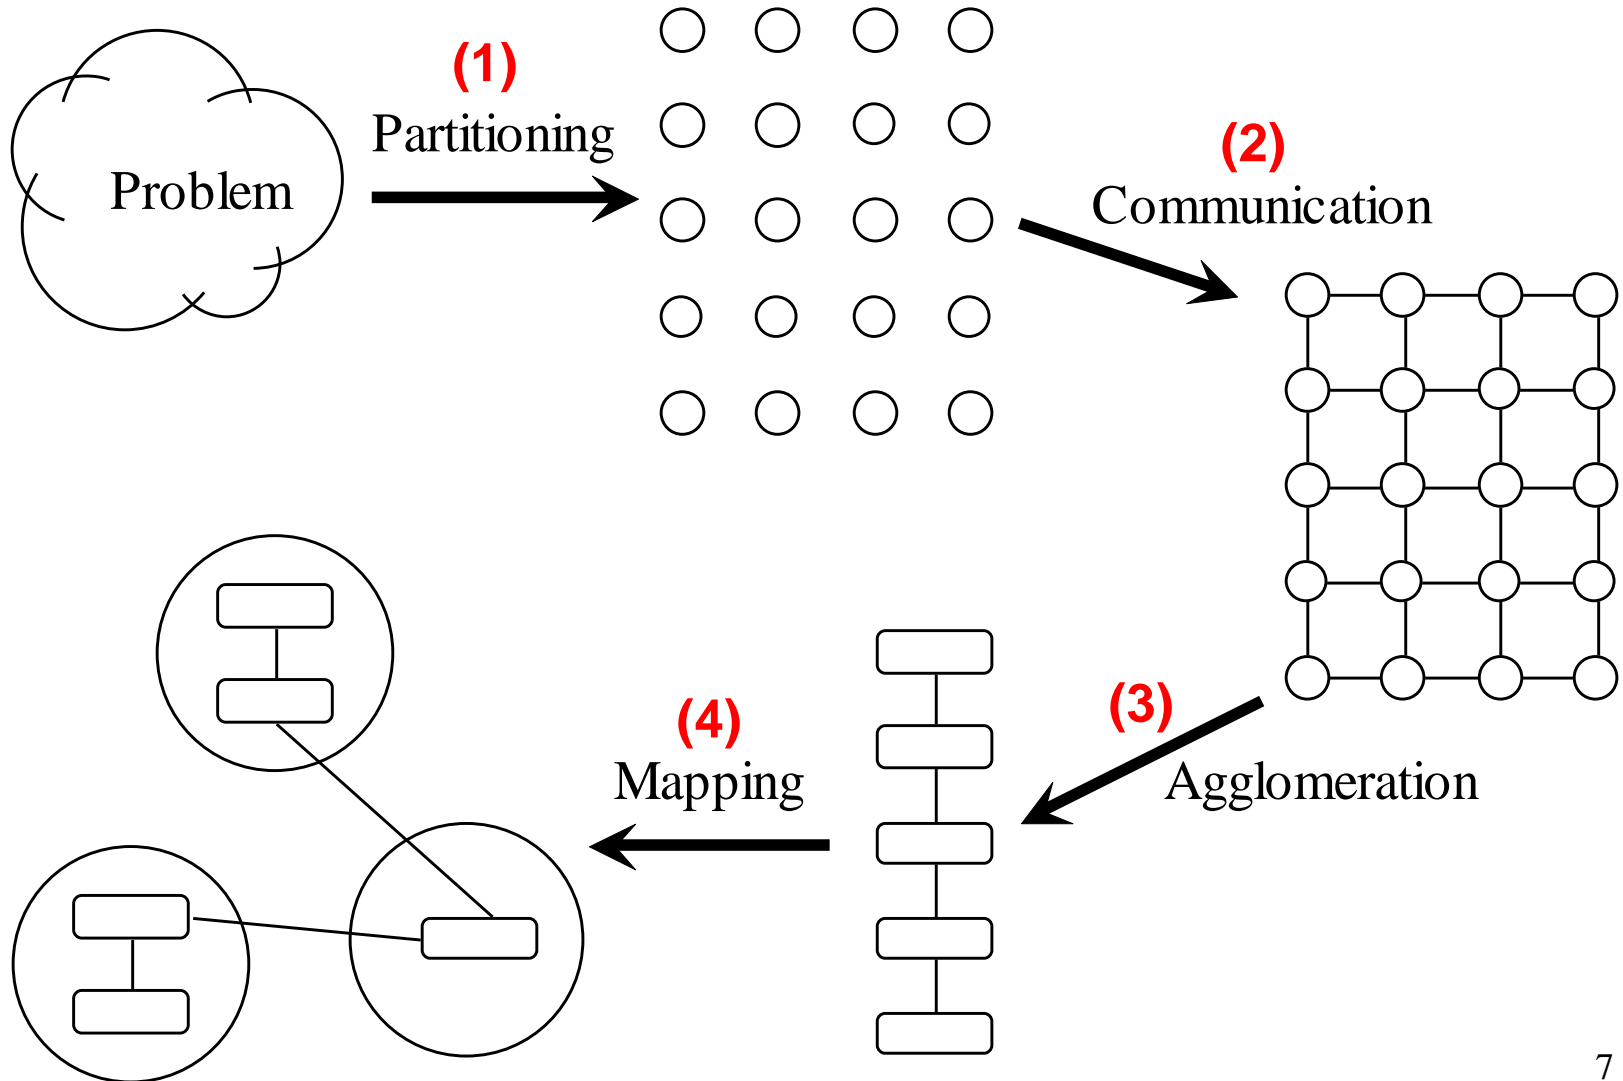
Task
= Program
+ Local Memory
+ I/O Ports

Task **Channel**

Foster's Design Methodology

- Ian Foster has proposed a 4-steps process for designing parallel algorithms
 1. Partitioning
 2. Communication
 3. Agglomeration
 4. Mapping

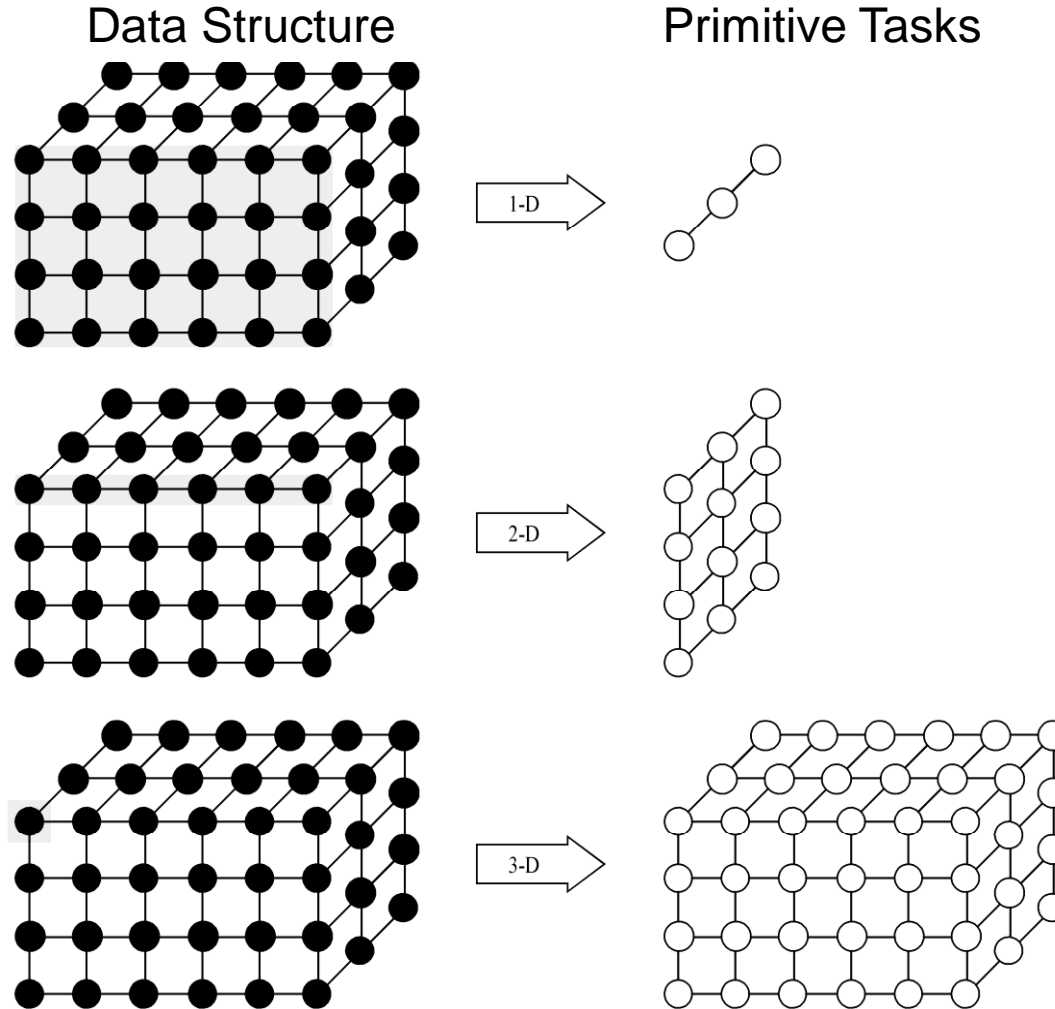
Foster's Methodology



Partitioning

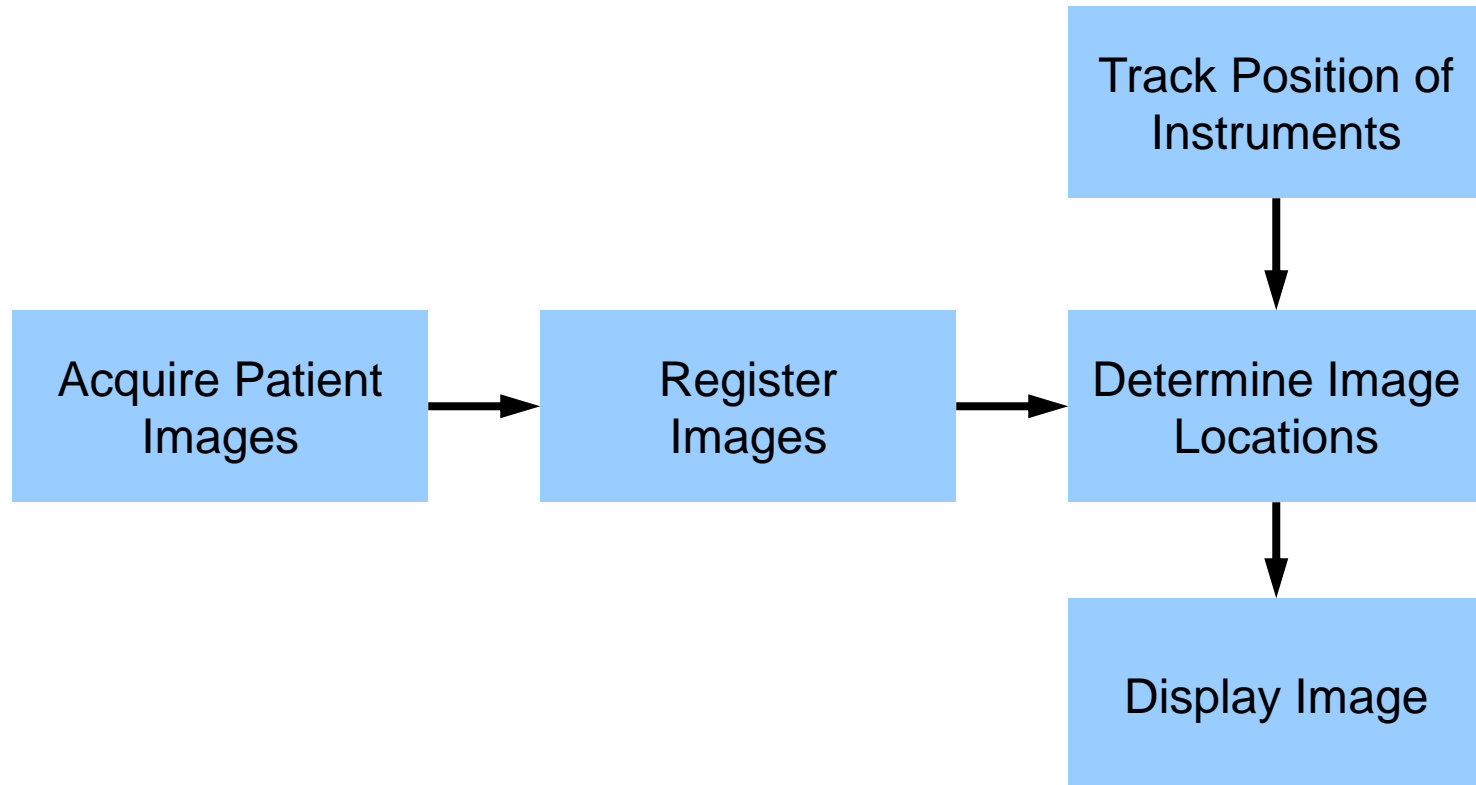
- Dividing **computation** and **data** into pieces.
 - Data-centric approach or computation-centric approach
- Domain decomposition
 - Divide **data** into pieces
 - Determine how to associate computations with the data
- Functional decomposition
 - Divide **computation** into pieces (e.g. pipelining)
 - Determine how to associate data with the computations

Example Domain Decompositions



Three **domain decompositions** of a 3D matrix

Example Functional Decomposition



Functional decomposition of a system supporting
interactive image-guided surgery

Partitioning Checklist

- **More tasks:** At least $10\times$ more primitive tasks than processors in target parallel computer.
- **Less redundancy:** Minimize redundant computations and redundant data structure storage.
- **Same size:** Primitive tasks are roughly the same size.
- **Scalability:** The number of tasks is an increasing function of the problem size.

Communication

- Determine values passed among tasks
- Local communication
 - Task needs values from a small number of other tasks
 - Create channels illustrating data flow
- Global communication
 - Significant number of tasks contribute data to perform a computation
 - Don't create channels for them early in design

Communication Checklist



- **Equal traffic:** Communication operations balanced among tasks
- **Less traffic:** Each task communicates with only small group of neighbors
- **Talking:** Tasks can perform communications concurrently
- **Working:** Tasks can perform computations concurrently

Agglomeration

- Grouping tasks into larger tasks
 - Improve performance (a.)
 - Simplify programming (b.)
- Goals
 - Lower communication overhead (a.)
 - Maintain scalability of the parallel design (a.)
 - Reduce software engineering cost (b.)
- In MPI programming, goal often to create one agglomerated task per processor

Agglomeration Can Improve Performance

- **Increasing the locality:** Eliminate communication between primitive tasks agglomerated into consolidated task
- **Reducing channels:** Combine groups of sending and receiving tasks



Agglomeration Checklist


Tuning partitioning and communication.

- Locality of parallel algorithm has increased
- Replicated computations take less time than communications they replace
- Data replication doesn't affect scalability
- Agglomerated tasks have similar computational and communications costs
- Number of tasks increases with problem size
- Number of tasks suitable for likely target systems
- **Trade-off** between agglomeration and code modifications costs is reasonable

Mapping


- Process of assigning tasks to processors
- Centralized multiprocessor
 - mapping done by **operating system**
- Distributed memory system
 - mapping done by **user** (our target)
- **Conflicting** goals of mapping
 - Maximize processor utilization
 - Minimize inter-processor communication

Mapping Example



(a)

Task/channel graph



(b)

Mapping of tasks to three processors

Optimal Mapping

- Finding optimal mapping is NP-hard
 - There are no known polynomial-time algorithms to map tasks to processors to minimize the execution time.
 - NP-hard
 - NP-hard, NIST, <http://www.nist.gov/dads/HTML/nphard.html>
 - NP-hard, Wikipedia, <http://en.wikipedia.org/wiki/NP-hard>
- Must rely on heuristics

Decision Tree of Mapping Strategy (1)

- Static number of tasks
 - Structured communication
 - Constant computation time per task
 - **Agglomerate tasks to minimize communication**
 - **Create one task per processor**
 - Variable computation time per task
 - **Cyclically map tasks to processors**
 - Unstructured communication
 - **Use a static load balancing algorithm**
- Dynamic number of tasks

Decision Tree of Mapping Strategy (2)

- Static number of tasks
- Dynamic number of tasks
 - Frequent communications between tasks
 - **Use a dynamic load balancing algorithm**
 - Many short-lived tasks
 - **Use a run-time task-scheduling algorithm**


Mapping Checklist

- **1 vs. M:** Considered designs based on one task per processor and multiple tasks per processor
- **S vs. D:** Evaluated static and dynamic task allocation
- If dynamic task allocation chosen, task allocator is not a **bottleneck** to performance
- If static task allocation chosen, ratio of tasks to processors is at least **10:1**

Case Studies


- Boundary value problem
- Finding the maximum
- The n-body problem
- Adding data input

Boundary Value Problem




A thin rod is suspended between two ice baths.
Then ends of the rod are in contact with the icewater.
The rod is surrounded by a thick blanket of insulation.
We can use a **partial differential equation** to model the temperature at any point on the rod as a function of time

Rod Cools as Time Progresses



The **finite difference method** finds the temperature at a fixed number of points in the rod at certain time intervals. **Decreasing** the size of the steps in space and time can lead to more accurate solutions.

Finite Difference Approximation



$$u = 100\sin(\pi x)$$

$$u_{i,j+1} = ru_{i-1,j} + (1 - 2r)u_{i,j} + ru_{i+1,j}$$

$$r = \frac{k}{h^2}$$

Data structure used in a finite difference approximation to the rod-cooling problem. Every point $u_{i,j}$ represents a matrix element containing the temperature at position i on the rod at time j . At each end of the rod the temperature is always 0. At time 0 the temperature at point x is $100\sin(\pi x)$.


1. Partitioning

- One data item per grid point – **easy**
- Associate one primitive task with each grid point
- Two-dimensional domain decomposition


2. Communication

- Identify communication pattern between primitive task
- Each interior primitive task has three incoming and three outgoing channels

$$u_{i,j+1} = ru_{i-1,j} + (1-2r)u_{i,j} + ru_{i+1,j}$$




3. Agglomeration and 4. Mapping



$$(a) \quad u_{i,j+1} = ru_{i-1,j} + (1-2r)u_{i,j} + ru_{i+1,j}$$



(b) Element i for all time steps.



(c) Over all time steps.

Agglomeration



Sequential Execution Time

- The rod has been divided into n pieces of size h .

χ – time to update element $u_{i,j+1} (= ru_{i-1,j} + (1-2r)u_{i,j} + ru_{i+1,j})$

n – number of elements

m – number of iterations

- Sequential execution time: $m(n-1)\chi$

Parallel Execution Time

- If each processor is responsible for an equal-sized portion of the rod's elements.

p – number of processors

λ – message latency

- The computation time for each iteration:

$$\chi \left\lceil \frac{n-1}{p} \right\rceil$$

- Necessary communication time: 2λ

- Parallel execution time:

$$m \left(\chi \left\lceil \frac{n-1}{p} \right\rceil + 2\lambda \right)$$

Finding the Maximum Error

- The error between the computed solution x and correction solution is $|(x - c) / c|$.
- Enhance previous parallel algorithm to find the maximum error.

Computed	0.15	0.16	0.16	0.19
Correct	0.15	0.16	0.17	0.18
Error (%)	0.00%	0.00%	6.25%	5.26%



6.25%

Reduction


- Given
 - a set of n values $a_0, a_1, a_2, \dots, a_{n-1}$
 - an associative operator \oplus
- Reduction is the process of computing $a_0 \oplus a_1 \oplus a_2 \oplus \dots \oplus a_{n-1}$
 - Examples: add, multiply, AND, OR, **maximum**, **minimum**
- Reduction requires exactly $n-1$ operations, it has **$\Theta(n)$** time complexity on a sequential computer.
 - How to perform a reduction on parallel computer quickly?

Parallel Reduction Evolution (1)

- Partitioning
 - Divide it into n pieces, one task per piece.
- Communication

χ – time to perform an addition


λ – message latency



Only one task receives
all other $n-1$ results.

$$\text{Time complexity} \\ = (n-1)(\lambda + \chi)$$

Parallel Reduction Evolution (2)



Two tasks work together.


$$2\left(\frac{n}{2} - 1\right) + 1 = n - 1$$

Time complexity

$$= \left(\frac{n}{2} - 1\right)(\lambda + \chi) + (\lambda + \chi)$$

$$= \frac{n}{2}(\lambda + \chi)$$

Parallel Reduction Evolution (3)




Four task cooperate. $4\left(\frac{n}{4} - 1\right) + 3 = n - 1$

Time complexity

$$= \left(\frac{n}{4} - 1\right)(\lambda + \chi) + 2(\lambda + \chi)$$

$$= \left(\frac{n}{4} + 1\right)(\lambda + \chi)$$

Continue ... Binomial Trees




Subgraph of Hypercube


$$n = 2^k$$
$$k = \log n$$

Binomial trees with 1, 2, 4, and 8 nodes.


Finding Global Sum in logarithmic time




Finding Global Sum




Finding Global Sum



Finding Global Sum




Finding Global Sum



Binomial Tree


Agglomeration and Mapping



n tasks mapped to
p processors

16 tasks are mapped to 4 processors.

Agglomerate Primitive Tasks



4 tasks on each processor are agglomerated
into a single task

n/p primitive tasks
with **1** value



1 primitive tasks
with n/p values

Analysis


χ – time to perform the binary operation

λ – message latency



- All tasks performance concurrently: $\left(\left\lceil \frac{n}{p} \right\rceil - 1 \right) \chi$
- A reduction of p values distributed among p tasks can be preformed in $\lceil \log p \rceil$ communication steps.
- Each reduction stop requires time: $\lambda + \chi$
- Overall execution time: $\left(\left\lceil \frac{n}{p} \right\rceil - 1 \right) \chi + \lceil \log p \rceil (\lambda + \chi)$

The n-Body Problem

- Newtonian n-body simulation.






Straightforward sequential algorithms
Time complexity: $\Theta(n^2)$ per iteration.



The n-Body Problem


Straightforward sequential algorithms
Time complexity: $\Theta(n^2)$ per iteration.







Partitioning

- Domain partitioning
- Assume one task per particle
- Task
 - particle's position
 - velocity vector
- Iteration
 - Get positions of all other particles
 - Compute new position, velocity


Gather



All-gather




Complete Graph for All-gather



Set up a channel between every pair of tasks.

Hypercube for All-gather



Each task have only **$\log p$** outgoing channels
and **$\log p$** incoming channels

Analysis

β – bandwidth of each channel

- Communication time


- Complete graph

$$(p-1)\left(\lambda + \frac{n/p}{\beta}\right) = (p-1)\lambda + \frac{n(p-1)}{\beta p}$$

- Hypercube

$$\sum_{i=1}^{\log p} \left(\lambda + \frac{2^{i-1} n/p}{\beta} \right) = \lambda \log p + \frac{n(p-1)}{\beta p}$$

- Overall time complexity




$$\lambda \log p + \frac{n(p-1)}{\beta p} + \chi\left(\frac{n}{p}\right)(n-1)$$

Error in textbook

Adding Data Input

- Augment the task/channel graph for the n-body problem




I/O Time Complexity :

$$\lambda_{io} + \frac{n}{\beta_{io}} \left\{ \text{for } n \text{ data elements (1D)} \right\}$$


$$\lambda_{io} + \frac{4n}{\beta_{io}} \left\{ \begin{array}{l} \text{for positions (2D)} \\ \text{and velocities (2D) of } n \text{ particles} \end{array} \right\}$$

Task 0 is responsible for I/O.

Scatter



Scatter in $\log p$ Steps



Communication

- Time for scatter the particles (sequential)
 - Send $p-1$ message, each of length $4n/p$

$$(p-1)\left(\lambda + \frac{4n}{p\beta}\right) = (p-1)\lambda + \frac{4n(p-1)}{p\beta}$$

- Time for scatter the particles (in $\log p$ steps)

$$\sum_{i=1}^{\log p} \left(\lambda + \frac{4n}{2^i \beta} \right) = \lambda \log p + \frac{4n(p-1)}{p\beta}$$

Error in textbook

- The “ $\log p$ steps” algorithm is better.

Analysis

- Overall execution time for m iterations:

$$2\left(\lambda_{io} + \frac{4n}{\beta_{io}}\right)$$

Input and Output

$$+ 2\left(\lambda \log p + \frac{4n(p-1)}{p\beta}\right)$$

Log p steps
scattering

$$+ m\left(\lambda \log p + \frac{2n(p-1)}{p\beta} + \chi \left\lceil \frac{n}{p} \right\rceil (n-1)\right)$$

All-gather and
computation

Summary: Task/channel Model

- Parallel computation
 - Set of tasks
 - Interactions through channels
- Good designs
 - Maximize local computations
 - Minimize communications
 - Scale up

Summary: Design Steps

- Partition computation
- Agglomerate tasks
- Map tasks to processors
- Goals
 - Maximize processor utilization
 - Minimize inter-processor communication

Summary: Fundamental Algorithms

- Reduction
- Gather and scatter
- All-gather


Exercise

- 3.11
- 3.13
- 3.18, 3.19

Discussion

- Real world
 - Processor
 - Single Core
 - Hyper-threading, multi-threading
 - <http://www.intel.com/personal/desktop/dualcore/demo/popup/demo.htm>
 - Dual Core
 - Programming
 - Multi-thread (share memory space)
 - Multi-process (no share memory space)
 - extend to MPI process.

Parallelization Policy



1-to-1 mapping?

