

Parallel Programming in C with MPI and OpenMP

Michael J. Quinn



Chapter 11

Matrix Multiplication




Outline

- Sequential algorithms
 - Iterative, row-oriented
 - Recursive, block-oriented
- Parallel algorithms
 - Rowwise block striped decomposition
 - Cannon's algorithm

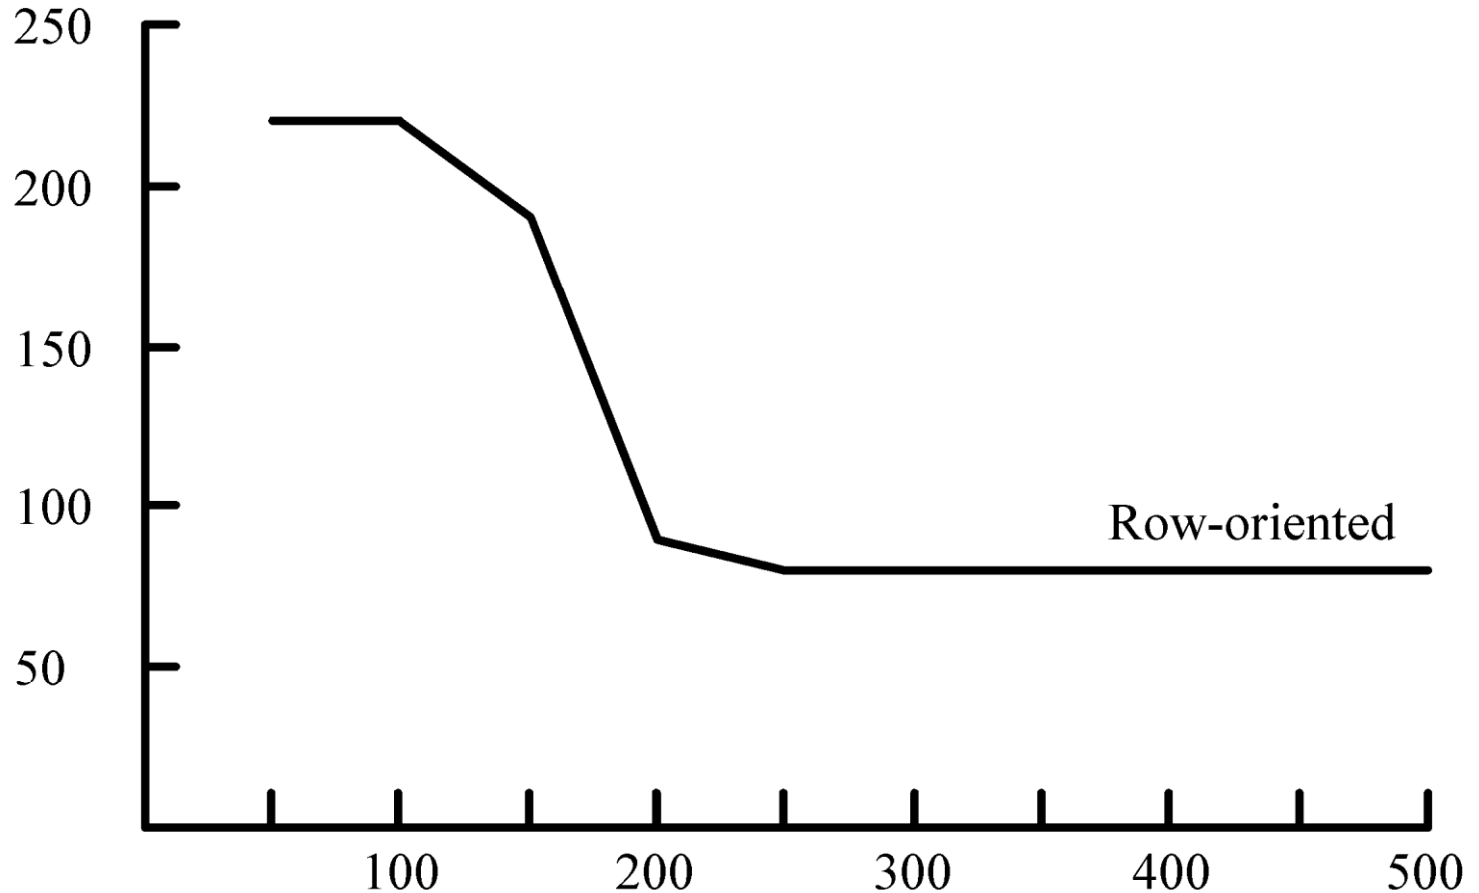
Iterative, Row-oriented Algorithm

- Series of inner product (dot product) operations



Performance as n Increases

Megaflops




Matrix Size

Row-oriented

Reason:


Matrix B Gets Too Big for Cache

- Computing a row of C requires accessing every element of B



Block Matrix Multiplication

- Replace scalar multiplication with matrix multiplication
- Replace scalar addition with matrix addition



Recurse Until B Small Enough

A


1	2	

B

1		
2		

C

Comparing Sequential Performance



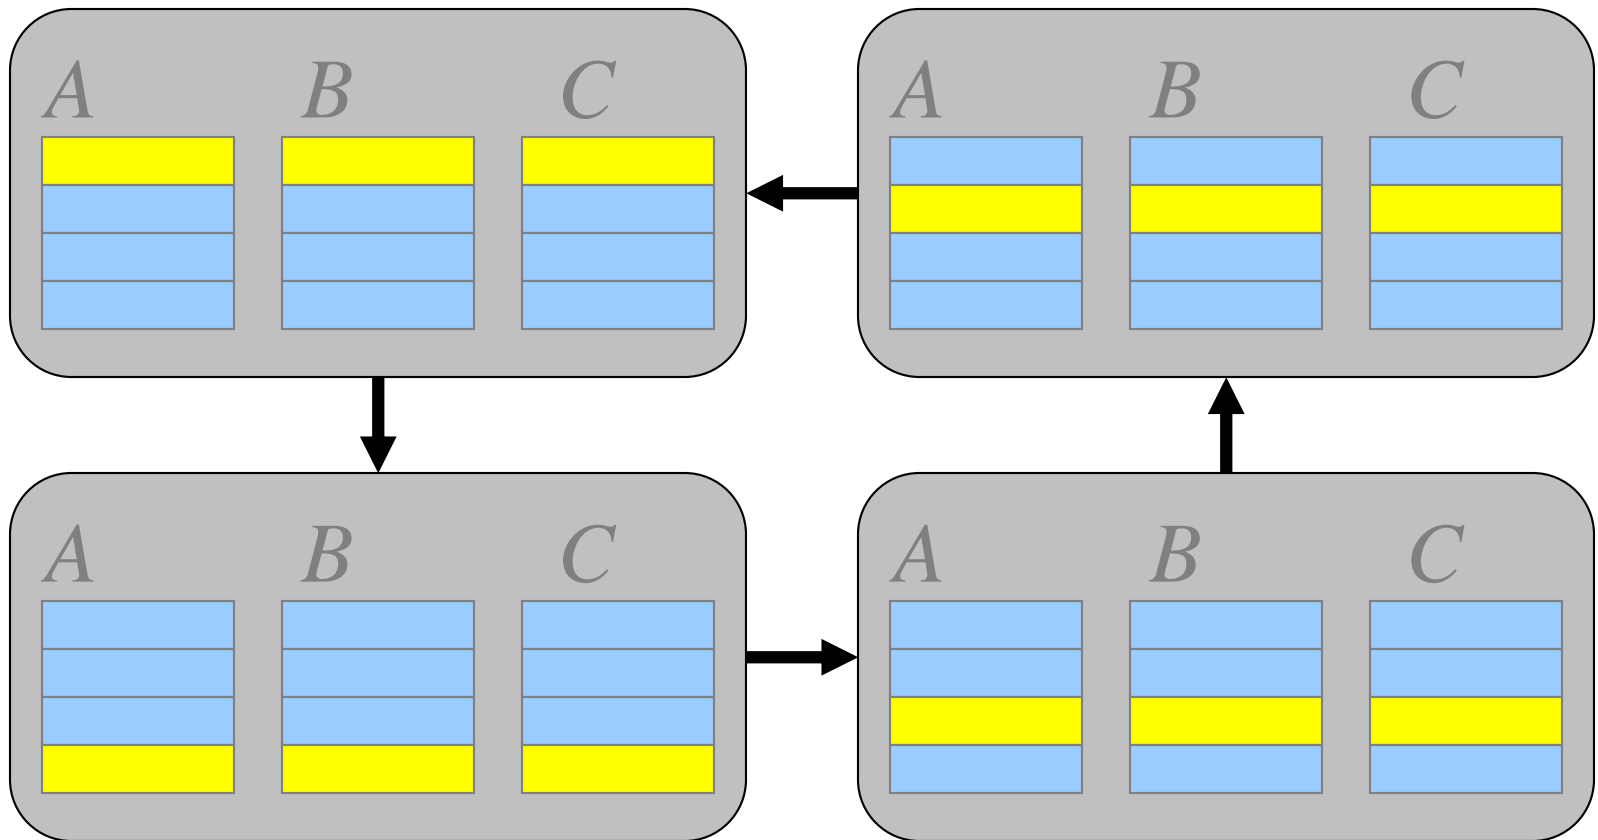
First Parallel Algorithm

- Partitioning
 - Divide matrices into rows
 - Each primitive task has corresponding rows of three matrices
- Communication
 - Each task must eventually see every row of B
 - Organize tasks into a ring

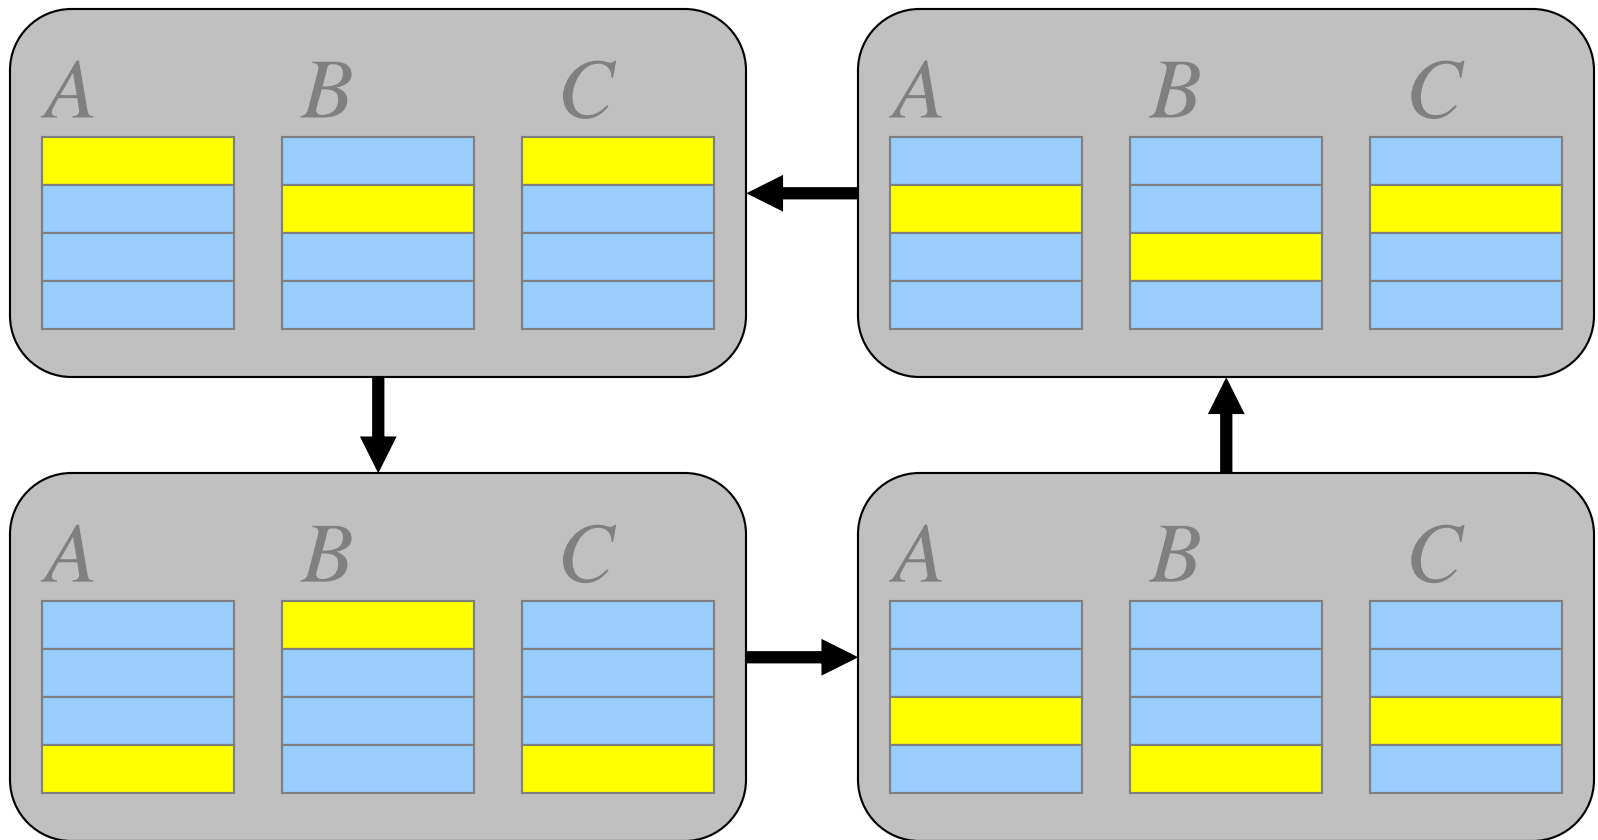
First Parallel Algorithm (cont.)

- Agglomeration and mapping
 - Fixed number of tasks, each requiring same amount of computation
 - Regular communication among tasks
 - Strategy: Assign each process a contiguous group of rows


Communication of B (1)




Communication of B (2)



Communication of B (3)



Communication of B (4)



Complexity Analysis

- Algorithm has p iterations
- During each iteration a process multiplies $(n / p) \times (n / p)$ block of A by $(n / p) \times n$ block of B : $\Theta(n^3 / p^2)$
- Total computation time: $\Theta(n^3 / p)$
- Each process ends up passing $(p - 1)n^2/p = \Theta(n^2)$ elements of B

Isoefficiency Analysis

- Sequential algorithm: $\Theta(n^3)$
- Parallel overhead: $\Theta(pn^2)$
Isoefficiency relation: $n^3 \geq Cpn^2 \Rightarrow n \geq Cp$

$$M(Cp)/p = C^2 p^2 / p = C^2 p$$

- This system does not have good scalability



Weakness of Algorithm 1

- Blocks of B being manipulated have p times more columns than rows
- Each process must access every element of matrix B
- Ratio of computations per communication is poor: only $2n / p$

Parallel Algorithm 2 (Cannon's Algorithm)

- Associate a primitive task with each matrix element
- Agglomerate tasks responsible for a square (or nearly square) block of C
- Computation-to-communication ratio rises to n / \sqrt{p}

Elements of A and B Needed to Compute a Process's Portion of C



Blocks Must Be Aligned

$A_{0,0}$ $B_{0,0}$	$A_{0,1}$ $B_{0,1}$	$A_{0,2}$ $B_{0,2}$	$A_{0,3}$ $B_{0,3}$
------------------------	------------------------	------------------------	------------------------

$A_{0,0}$ $B_{0,0}$	$A_{0,1}$ $B_{1,1}$	$A_{0,2}$ $B_{2,2}$	$A_{0,3}$ $B_{3,3}$
------------------------	------------------------	------------------------	------------------------

$A_{1,0}$ $B_{1,0}$	$A_{1,1}$ $B_{1,1}$	$A_{1,2}$ $B_{1,2}$	$A_{1,3}$ $B_{1,3}$
------------------------	------------------------	------------------------	------------------------

$A_{1,1}$ $B_{1,0}$	$A_{1,2}$ $B_{2,1}$	$A_{1,3}$ $B_{3,2}$	$A_{1,0}$ $B_{0,3}$
------------------------	------------------------	------------------------	------------------------

$A_{2,0}$ $B_{2,0}$	$A_{2,1}$ $B_{2,1}$	$A_{2,2}$ $B_{2,2}$	$A_{2,3}$ $B_{2,3}$
------------------------	------------------------	------------------------	------------------------

$A_{2,2}$ $B_{2,0}$	$A_{2,3}$ $B_{3,1}$	$A_{2,0}$ $B_{0,2}$	$A_{2,1}$ $B_{1,3}$
------------------------	------------------------	------------------------	------------------------

$A_{3,0}$ $B_{3,0}$	$A_{3,1}$ $B_{3,1}$	$A_{3,2}$ $B_{3,2}$	$A_{3,3}$ $B_{3,3}$
------------------------	------------------------	------------------------	------------------------


$A_{3,3}$ $B_{3,0}$	$A_{3,0}$ $B_{0,1}$	$A_{3,1}$ $B_{1,2}$	$A_{3,2}$ $B_{2,3}$
------------------------	------------------------	------------------------	------------------------

Before


After

Blocks Need to Be Aligned

- Each triangle represents a matrix block
- Only same-color triangles should be multiplied




Rearrange Blocks



Block A_{ij} cycles
left i positions


Block B_{ij} cycles
up j positions

Consider Process $P_{1,2}$




Step 1

Consider Process $P_{1,2}$




Step 2

Consider Process $P_{1,2}$



Step 3

Consider Process $P_{1,2}$



Step 4

Complexity Analysis

- Algorithm has \sqrt{p} iterations
- During each iteration process multiplies two $(n / \sqrt{p}) \times (n / \sqrt{p})$ matrices: $\Theta(n^3 / p^{3/2})$
- Computational complexity: $\Theta(n^3 / p)$
- During each iteration process sends and receives two blocks of size $(n / \sqrt{p}) \times (n / \sqrt{p})$
- Communication complexity: $\Theta(n^2 / \sqrt{p})$

Isoefficiency Analysis

- Sequential algorithm: $\Theta(n^3)$
- Parallel overhead: $\Theta(\sqrt{pn^2})$

Isoefficiency relation: $n^3 \geq C \sqrt{pn^2} \Rightarrow n \geq C \sqrt{p}$

$$M(C\sqrt{p})/p = C^2 p/p = C^2$$

- This system is highly scalable

Summary

- Considered two sequential algorithms
 - Iterative, row-oriented algorithm
 - Recursive, block-oriented algorithm
 - Second has better cache hit rate as n increases
- Developed two parallel algorithms
 - First based on rowwise block striped decomposition
 - Second based on checkerboard block decomposition
 - Second algorithm is scalable, while first is not