



## Data Structures: Lecture 14

Herbert H. Chang, Ph.D.

張恆華

Eng Sci Ocean Eng

National Taiwan University

*Linear-Time Sorting*

H. Chang © NTU

## Bucket Sort

- Works well when keys are in small range, e.g., from 0 to  $q-1$  and the number of items  $n$  is larger than, or nearly as large as,  $q$ .  $\rightarrow$  i.e. when  $q \in O(n)$ .
- Allocate an array of  $q$  queues, numbered from 0 to  $q-1$ . The queues are called “buckets”.
- Walk through the list of input items and enqueue each item: an item with key  $i$  goes into queue  $i$ .
- Input: each item has a numerical key and an associated value.

2


H. Chang © NTU

# Illustration

- Input: each item has a numerical key and an associated value.

6:a	7:b	3:c	0:d	3:e	1:f	5:g	0:h	3:i	7:j	n
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	---

- Queues:



- Concatenate queues together in order. (→)
- This data structure is exactly like a hash table, no compression fcn.

3

H. Chang © NTU

# Running Time

- Runs in  $\Theta(q+n)$  time (best- and worst-case).
  - $\Theta(q)$  time to initialize the buckets in the beginning and to concatenate them together in the end.
  - $\Theta(n)$  time to put all items in their buckets.
- If  $q \in O(n)$ , that is, the number of possible keys isn't much larger than the number of items we're sorting, then total is  $\Theta(n)$  time.

4

H. Chang © NTU

## Stable Sorting

- Bucket sort is said to be stable.
- A sort is stable if items with equal keys come out in the same order they went in.
- Insertion, selection, mergesort are easily made stable.
- Linked list quicksort can too; but array version is not.
- Heapsort is never stable.
- Bucket sort is ONLY appropriate when keys are distributed in a small range; i.e.  $q$  is in  $O(n)$ .

5

H. Chang © NTU

## Counting Sort

- If the items are naked keys with no associated values, bucket sort can be simplified to become “counting sort”.
- In counting sort, use no queues.
- Merely keep a count of how many copies of each key.
- Suppose we sort 6 7 3 0 3 1 5 0 3 7:

	0	1	2	3	4	5	6	7
counts	2	1	0	3	0	1	1	2


- When finished, output: 0 0 1 3 3 3 5 6 7 7

6

H. Chang © NTU

## Counting Sort with Complete Items

- Having complete items (key plus associated value).
- The trick is to use the counts to find the right index to move each item to.
- Let  $x$  be an input array of objects with keys.



7

H. Chang © NTU

## Counting Sort with Complete Items

- Begin by counting the keys in  $x$
- ```
int lengthx = sizeof(x)/sizeof(x[0]);  
for (i=0; i<lengthx; i++) {  
    counts[x[i].key]++;  
}
```

|        | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
|--------|---|---|---|---|---|---|---|---|
| counts | 2 | 1 | 0 | 3 | 0 | 1 | 1 | 2 |

8

H. Chang © NTU



## Counting Sort with Complete Items

- Next, do a scan: so that counts[i] contains the number of keys “less than” i.

|        | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
|--------|---|---|---|---|---|---|---|---|
| counts | 0 | 2 | 3 | 3 | 6 | 6 | 7 | 8 |

```
total = 0;
lengthc = sizeof(counts)/sizeof(counts[0]);
for (j=0; j<lengthc; j++) {
    c = counts[j];
    counts[j] = total;
    total = total + c;
}
```

9

H. Chang © NTU


## Counting Sort with Complete Items

- Let y be the output array to put the sorted objects.
- counts[i] tells us the first index of y to put items with key i.
- Walk through the array x and copy each item to its final position in y.
- When you copy an item with key k, increment counts[k] to make sure that the next item with key k goes into the next slot.

10

H. Chang © NTU

## Counting Sort with Complete Items



11

H. Chang © NTU

## Running Time


- Bucket sort and counting sort both take  $O(q+n)$  time. If  $q \in O(n)$ , they take  $O(n)$  time.
- If you're sorting an array, counting sort takes less memory than bucket sort.
- If you're sorting a linked list, bucket sort more natural.
- What if  $q \gg n$ ?

12

H. Chang © NTU

# Radix Sort

- Suppose we want to sort 1,000 items in the range from 0 to 99,999,999.
- Instead of providing 100 million buckets, let's provide  $q = 10$  buckets and sort on the first digit only.
- We use bucket sort or counting sort, treating each item as if its key is the first digit of its true key.



13

H. Chang © NTU

# Radix Sort

- We could sort queue recursively on second digit, then on sorted third digit, and so on.
- Unfortunately, this tends to break the set of input items into too many small subsets, each of which will be sorted relatively inefficiently.
- Clever idea: we'll keep all numbers in one pile throughout sort, sort on last digit first, then next-to-last, up to the most significant digit.
- Why this works? Because bucket sort and counting sort are stable. Once we sort on the last digit, 55,555,552 and 55,555,58 remain sorted.

14

H. Chang © NTU

## Example with three-digit numbers

- Three-digit numbers:

Sort on 1s: 720 450 771 822 332 504 925 5 955 825 777 858 28 829

Sort on 10s: 504 5 720 822 925 825 28 829 332 450 955 858 771 777

Sort on 100s: 5 28 332 450 504 720 771 777 822 825 829 858 925 955

- It would likely be faster if we sort on 2 digits at a time (a radix of  $q=100$ ) or even 3 (a radix of  $q=1,000$ ).
- On computers, more natural to choose power-of-two radix like  $q=256$ . (1-byte at a time)
- Note that  $q = \text{\# of buckets} = \text{radix of digit}$  we use as a sort key during one pass of bucket or counting sort.

15

H. Chang © NTU

## Number of passes

- How many passes must we perform?
- Each pass inspects  $\log_2 q$  bits of each key.
- $b=32$ -bit int:  $q=256$



- If all the keys can be represented in  $b$  bits, the number of passes is  $\left\lceil \frac{b}{\log_2 q} \right\rceil$ .
- Running time of radix sort is  $O((n+q) \left\lceil \frac{b}{\log_2 q} \right\rceil)$ .

16

H. Chang © NTU



## Running time

- How to choose  $q$ ?  
Choose  $q \in O(n)$ , so each pass takes linear time.  
Make  $q$  large enough  $\rightarrow$  # of passes small.  
Let's choose  $q \approx n \rightarrow$   
Radix sort takes  $O(n + \frac{b}{\log n} n)$  time.
- For keys = ints:  $b$  is a constant; radix sort takes linear time.
- Practical choice: make  $q$  be  $n$  rounded down to the next power of two.
- To keep memory use low: make  $q \approx \sqrt{n}$ , rounded to the nearest power of two.